


Lect 19 – Data Aggregation

Rob Capra

INLS 490-172

Data files for today

- Sakai → Resources → Lectures → lect18_data.zip



MLB stats from last lecture

```
In [25]: !type mlb.csv
team,league,wins,losses,rs,ra
yankees,al,6,6,46,52
nationals,nl,7,5,60,50
cardinals,nl,7,5,48,48
redsox,al,5,7,44,50
braves,nl,8,4,46,33
cubs,nl,4,8,47,55
tigers,al,6,4,40,39
```

```
In [26]: stats_df = pd.read_csv('mlb.csv')
```

```
In [27]: print stats_df
```

	team	league	wins	losses	rs	ra
0	yankees	al	6	6	46	52
1	nationals	nl	7	5	60	50
2	cardinals	nl	7	5	48	48
3	redsox	al	5	7	44	50
4	braves	nl	8	4	46	33
5	cubs	nl	4	8	47	55
6	tigers	al	6	4	40	39

teams-cities.csv

- Suppose we have another, related set of data that indicates the city and state for each team

```
In [28]: !type teams-cities.csv
team,city,state
yankees,new york,ny
nationals,washington,dc
cardinals,st. louis,mo
redsox,boston,ma
braves,atlanta,ga
cubs,chicago,il
tigers,detroit,mi
```

```
In [29]: city_df = pd.read_csv('teams-cities.csv')
```

```
In [30]: print city_df
```

	team	city	state
0	yankees	new york	ny
1	nationals	washington	dc
2	cardinals	st. louis	mo
3	redsox	boston	ma
4	braves	atlanta	ga
5	cubs	chicago	il
6	tigers	detroit	mi

merge

- Database-style join/merge on DataFrames.

```
In [31]: print stats_df
```

	team	league	wins	losses	rs	ra
0	yankees	al	6	6	46	52
1	nationals	nl	7	5	60	50
2	cardinals	nl	7	5	48	48
3	redsox	al	5	7	44	50
4	braves	nl	8	4	46	33
5	cubs	nl	4	8	47	55
6	tigers	al	6	4	40	39

```
In [32]: print city_df
```

	team	city	state
0	yankees	new york	ny
1	nationals	washington	dc
2	cardinals	st. louis	mo
3	redsox	boston	ma
4	braves	atlanta	ga
5	cubs	chicago	il
6	tigers	detroit	mi

```
In [33]: print pd.merge(stats_df, city_df)
```

	team	league	wins	losses	rs	ra	city	state
0	yankees	al	6	6	46	52	new york	ny
1	nationals	nl	7	5	60	50	washington	dc
2	cardinals	nl	7	5	48	48	st. louis	mo
3	redsox	al	5	7	44	50	boston	ma
4	braves	nl	8	4	46	33	atlanta	ga
5	cubs	nl	4	8	47	55	chicago	il
6	tigers	al	6	4	40	39	detroit	mi

Merge on index

```
In [5]: stats_df = pd.read_csv('mlb.csv', index_col='team')
```

```
In [6]: print stats_df
```

	league	wins	losses	rs	ra
team					
yankees	al	6	6	46	52
nationals	nl	7	5	60	50
cardinals	nl	7	5	48	48
redsox	al	5	7	44	50
braves	nl	8	4	46	33
cubs	nl	4	8	47	55
tigers	al	6	4	40	39

```
In [7]: city_df = pd.read_csv('teams-cities.csv', index_col='team')
```

```
In [8]: print city_df
```

	city	state
team		
yankees	new york	ny
nationals	washington	dc
cardinals	st. louis	mo
redsox	boston	ma
braves	atlanta	ga
cubs	chicago	il
tigers	detroit	mi

```
In [9]: print pd.merge(stats_df, city_df, left_index=True, right_index=True)
```

	league	wins	losses	rs	ra	city	state
team							
yankees	al	6	6	46	52	new york	ny
nationals	nl	7	5	60	50	washington	dc
cardinals	nl	7	5	48	48	st. louis	mo
redsox	al	5	7	44	50	boston	ma
braves	nl	8	4	46	33	atlanta	ga
cubs	nl	4	8	47	55	chicago	il
tigers	al	6	4	40	39	detroit	mi

Merge on index and column

```
In [11]: print stats_df
```

	league	wins	losses	rs	ra
team					
yankees	al	6	6	46	52
nationals	nl	7	5	60	50
cardinals	nl	7	5	48	48
redsox	al	5	7	44	50
braves	nl	8	4	46	33
cubs	nl	4	8	47	55
tigers	al	6	4	40	39

```
In [12]: city_df = pd.read_csv('teams-cities.csv')
```

```
In [13]: print city_df
```

	team	city	state
0	yankees	new york	ny
1	nationals	washington	dc
2	cardinals	st. louis	mo
3	redsox	boston	ma
4	braves	atlanta	ga
5	cubs	chicago	il
6	tigers	detroit	mi

```
In [14]: print pd.merge(stats_df, city_df, left_index=True, right_on='team')
```

	league	wins	losses	rs	ra	team	city	state
0	al	6	6	46	52	yankees	new york	ny
1	nl	7	5	60	50	nationals	washington	dc
2	nl	7	5	48	48	cardinals	st. louis	mo
3	al	5	7	44	50	redsox	boston	ma
4	nl	8	4	46	33	braves	atlanta	ga
5	nl	4	8	47	55	cubs	chicago	il
6	al	6	4	40	39	tigers	detroit	mi

Using merge to filter

```
In [16]: print stats_df
```

	league	wins	losses	rs	ra
team					
yankees	al	6	6	46	52
nationals	nl	7	5	60	50
cardinals	nl	7	5	48	48
redsox	al	5	7	44	50
braves	nl	8	4	46	33
cubs	nl	4	8	47	55
tigers	al	6	4	40	39

```
In [17]: z = city_df[:3]
```

```
In [18]: print z
```

	team	city	state
0	yankees	new york	ny
1	nationals	washington	dc
2	cardinals	st. louis	mo

```
In [19]: print pd.merge(stats_df, z, left_index=True, right_on='team')
```

	league	wins	losses	rs	ra	team	city	state
0	al	6	6	46	52	yankees	new york	ny
1	nl	7	5	60	50	nationals	washington	dc
2	nl	7	5	48	48	cardinals	st. louis	mo

Recall – Read CSV + Hirearchical Index

- Remember that we can read in a CSV file
- And create a hierarchical index

```
In [44]: !type courses2.csv
course,sem,enroll,assign
inls101,f12,12,3
inls161,f12,18,4
inls382,f12,15,4
inls101,f13,17,4
inls161,f13,19,3
inls382,f13,21,5
```

```
In [45]: df = pd.read_csv('courses2.csv', index_col=['sem', 'course'])
```

```
In [46]: print df
```

		enroll	assign
sem	course		
f12	inls101	12	3
	inls161	18	4
	inls382	15	4
f13	inls101	17	4
	inls161	19	3
	inls382	21	5

Recall – Summary Stats

- Many summary stats functions have a *level* option that can be used with a hierarchical index

```
In [51]: print df
          enroll  assign
sem course
f12 inls101      12      3
     inls161      18      4
     inls382      15      4
f13 inls101      17      4
     inls161      19      3
     inls382      21      5
```

```
In [52]: print df.sum(level=0)
          enroll  assign
sem
f12         45      11
f13         57      12
```

```
In [53]: print df.sum(level=1)
          enroll  assign
course
inls101      29      7
inls161      37      7
inls382      36      9
```

```
In [54]: print
df['enroll'].sum(level=0)
sem
f12      45
f13      57
dtype: int64
```

```
In [55]: print
df.sum(level=0)['enroll']
sem
f12      45
f13      57
Name: enroll, dtype: int64
```

Aggregation – GroupBy

- As an alternative to creating and using a hierarchical index to do aggregation, we can use groupby.
- Group by uses a split-apply-combine process.

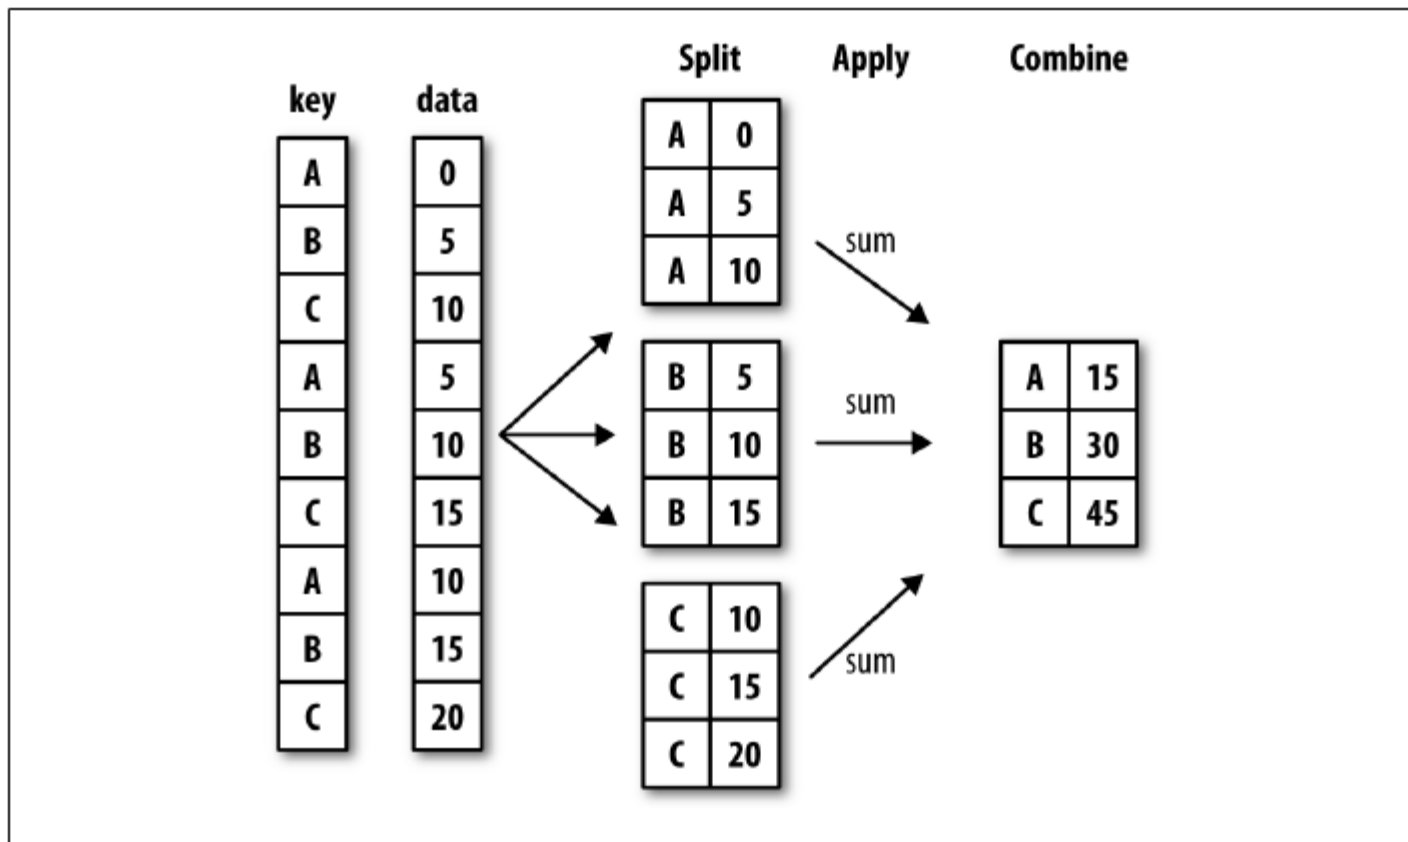


Figure 9-1. Illustration of a group aggregation

Aggregation – GroupBy

- In our previous examples, we created summary statistics by relying on a hierarchical index
- Using GroupBy, we will not need to have a hierarchical index.

Previous way – Hierarchical Index

```
In [62]: !type courses2.csv
course,sem,enroll,assign
inls101,f12,12,3
inls161,f12,18,4
inls382,f12,15,4
inls101,f13,17,4
inls161,f13,19,3
inls382,f13,21,5
```

```
In [63]: df =
pd.read_csv('courses2.csv',
index_col=['sem', 'course'])
```

```
In [64]: print df
           enroll  assign
sem course
f12 inls101      12       3
    inls161      18       4
    inls382      15       4
f13 inls101      17       4
    inls161      19       3
    inls382      21       5
```

New way – No hierarchical index

```
In [77]: df = pd.read_csv('courses2.csv')
```

```
In [78]: print df
   course  sem  enroll  assign
0  inls101  f12      12       3
1  inls161  f12      18       4
2  inls382  f12      15       4
3  inls101  f13      17       4
4  inls161  f13      19       3
5  inls382  f13      21       5
```

GroupBy Objects

- `.groupby()` on a DataFrame returns a GroupBy object
- GroupBy objects have methods such as `.sum()` and `.mean()`

```
In [84]: print df
   course sem enroll assign
0  inls101 f12     12      3
1  inls161 f12     18      4
2  inls382 f12     15      4
3  inls101 f13     17      4
4  inls161 f13     19      3
5  inls382 f13     21      5
```

```
In [85]: g = df['enroll'].groupby(df['sem'])
```

```
In [86]: print g
<pandas.core.groupby.SeriesGroupBy object at 0x0000000011422860>
```

```
In [87]: print g.sum()
sem
f12     45
f13     57
dtype: int64
```

```
In [88]: print g.mean()
sem
f12     15
f13     19
dtype: int64
```

```
In [95]: print type(g.sum())
<class 'pandas.core.series.Series'>
```

```
In [103]: print g.sum().index
Index([u'f12', u'f13'], dtype=object)
```

Creates a GroupBy object grouped by the semester. This object can be used later to do operations (such as sum and mean on the groups.

Also notice that `g.sum()` returns a Series with semester as its index

GroupBy

- If you groupby multiple columns, when you perform an operation such as `.sum()`, the result will have a hierarchical index

```
In [96]: print df
```

	course	sem	enroll	assign
0	inls101	f12	12	3
1	inls161	f12	18	4
2	inls382	f12	15	4
3	inls101	f13	17	4
4	inls161	f13	19	3
5	inls382	f13	21	5

```
In [97]: g = df['enroll'].groupby([df['sem'], df['course']])
```

```
In [98]: print g.sum()
```

sem	course	
f12	inls101	12
	inls161	18
	inls382	15
f13	inls101	17
	inls161	19
	inls382	21

```
dtype: int64
```

```
In [99]: print type(g.sum())
```

```
<class 'pandas.core.series.Series'>
```

```
In [100]: print g.sum().index
```

```
MultiIndex
```

```
[(u'f12', u'inls101'), (u'f12', u'inls161'), (u'f12', u'inls382'), (u'f13',  
u'inls101'), (u'f13', u'inls161'), (u'f13', u'inls382')]
```

Groupby + Unstack

- After doing a groupby with two columns, you may want to use unstack

```
In [111]: print df
   course sem enroll assign
0  inls101 f12      12      3
1  inls161 f12      18      4
2  inls382 f12      15      4
3  inls101 f13      17      4
4  inls161 f13      19      3
5  inls382 f13      21      5
```

```
In [112]: g = df['enroll'].groupby([df['sem'], df['course']])
```

```
In [113]: print g.sum()
```

```
sem  course
f12  inls101      12
      inls161      18
      inls382      15
f13  inls101      17
      inls161      19
      inls382      21
```

```
dtype: int64
```

```
In [114]: print g.sum().unstack()
```

```
course inls101 inls161 inls382
sem
f12           12         18         15
f13           17         19         21
```

```
In [115]: print g.sum().unstack()['inls382']['f13']
21
```

Groupby shorthand

- If the grouping information is in the same DataFrame as the data being aggregated, you can use a shorthand notation.

```
In [119]: print df
   course  sem  enroll  assign
0  inls101  f12      12       3
1  inls161  f12      18       4
2  inls382  f12      15       4
3  inls101  f13      17       4
4  inls161  f13      19       3
5  inls382  f13      21       5
```

```
In [120]: print df.groupby(df['sem']).sum()
   enroll  assign
sem
f12      45      11
f13      57      12
```

```
In [121]: print df.groupby('sem').sum()
   enroll  assign
sem
f12      45      11
f13      57      12
```


Watch out!

```
In [130]: print df
   course  sem  enroll  assign
0  inls101  f12      12       3
1  inls161  f12      18       4
2  inls382  f12      15       4
3  inls101  f13      17       4
4  inls161  f13      19       3
5  inls382  f13      21       5
```

```
In [131]: print df.groupby('sem').sum()
   enroll  assign
sem
f12      45      11
f13      57      12
```

```
In [132]: print df.groupby('sem').sum()['enroll']
sem
f12      45
f13      57
Name: enroll, dtype: int64
```

Line 133 does not work!
Why not?

```
In [133]: print df['enroll'].groupby('sem').sum() X
```

```
In [134]: print df['enroll'].groupby(df['sem']).sum()
sem
f12      45
f13      57
dtype: int64
```

```
In [146]: df.groupby('sem')['enroll'].sum()
Out[146]:
sem
f12      45
f13      57
Name: enroll, dtype: int64
```

Line 146 is another syntax that works

Using the results from groupby

```
In [136]: print df
```

	course	sem	enroll	assign
0	inls101	f12	12	3
1	inls161	f12	18	4
2	inls382	f12	15	4
3	inls101	f13	17	4
4	inls161	f13	19	3
5	inls382	f13	21	5

```
In [137]: z = df['enroll'].groupby(df['course']).sum()
```

```
In [138]: print z
```

```
course
inls101    29
inls161    37
inls382    36
dtype: int64
```

```
In [139]: z.sort(ascending=False)
```

```
In [140]: print z
```

```
course
inls161    37
inls382    36
inls101    29
dtype: int64
```

```
In [141]: type(z)
```

```
Out[141]: pandas.core.series.Series
```

```
In [143]: z[:2]
```

```
Out[143]:
```

```
course
inls161 37
inls382 36
dtype: int64
```

```
In [144]: for course, enroll in z[:2].iteritems():
...:     print course, enroll
...:
inls161 37
inls382 36
```

Manipulating DF, Series, Groupby

```
In [210]: print df
```

	course	sem	enroll	assign
0	inls101	f12	12	3
1	inls161	f12	18	4
2	inls382	f12	15	4
3	inls101	f13	17	4
4	inls161	f13	19	3
5	inls382	f13	21	5

```
In [211]: z =  
df.groupby('course')['enroll'].sum().order  
(ascending=False)[:2]
```

```
In [212]: print z  
course  
inls161    37  
inls382    36  
Name: enroll, dtype: int64
```

```
In [213]: zdf = DataFrame(z.values,  
index=z.index)
```

```
In [214]: print zdf  
0  
course  
inls161  37  
inls382  36
```

```
In [215]: zdf.columns = ['enroll']
```

```
In [216]: print zdf  
enroll  
course  
inls161    37  
inls382    36
```

```
In [218]: y = DataFrame(['Tools',  
'InfoSys'], index=['inls161', 'inls382'])
```

```
In [219]: print y  
0  
inls161    Tools  
inls382    InfoSys
```

```
In [220]: y.columns = ['coursename']
```

```
In [221]: print y  
coursename  
inls161    Tools  
inls382    InfoSys
```

```
In [222]: zdfy = zdf.join(y)
```

```
In [223]: print zdfy  
enroll coursename  
course  
inls161    37      Tools  
inls382    36      InfoSys
```

GroupBy Exercise

(not to turn in)

- Create a DataFrame using with the following data:
- Do NOT create a hierarchical index

- After creating the DF, use groupby to:
 1. Output a summary table of the total plays for each uid for each month (i.e. collapse artists)
 2. Output a summary table of the total plays for each artist for each month (i.e. collapse uids)
 3. Output a summary table of the total plays for each uid for each artists (i.e. collapse months)

		Aug	Sep	Nov
uid123	Bowie	12	15	26
uid123	Gaga	2	0	4
uid123	Spears	1	0	3
uid345	Bowie	3	0	4
uid345	Gaga	24	18	31
uid345	Spears	8	12	5
uid678	Bowie	6	3	0
uid678	Gaga	8	14	27
uid678	Spears	28	21	16